



Caruca v2

LLM-based specification mining for opaque shell commands

Master Idea Document

Version 1.1 — corrected 22 September 2026

SOURCE `ai_docs/prep/master_idea.md`

GENERATED September 22, 2026 at 4:34 PM

PAGES 8

CARUCA V2 — RESEARCH PROJECT

Table of contents

Master Idea Document

End Goal

Specific Problem

Stakeholders & Consumers

Evaluation Criteria & Success Definition

MVP Components (Flexible, provisional)

Key Usage Scenarios

Future / Stretch Directions (not committed)

Master Idea Document

Version 1.1 — corrected 22 September 2026. The 6,520 LOC figure was quoted bare as the size of the hand-written pipeline. About half of it is generated specification data; the hand-written figure is ~3,456. Corrected inline. Nothing else in this document changed.

End Goal

With `caruca_v2`, I want to determine, through rigorous instrumented comparison, whether an LLM-based approach can match or exceed Caruca v1's performance, correctness, and coverage, not just assume it can. I'll do this by building a naive single-prompt LLM baseline first, measuring it three-way against v1 and ground truth on cost, correctness, and percent-change, with a secure sandboxed execution environment for testing arbitrary, including destructive, command permutations. An agentic Claude Code pipeline rebuild follows later as separate work.

This works toward two objectives: near-term, proving that an LLM-based approach can perform Caruca-style specification mining at all and demonstrating its efficacy; longer-term, producing the evidence and documentation needed to revise and resubmit the Caruca white paper (not accepted in its original submission), incorporating v2's findings and a clear account of what changed between v1 and v2 and why.

Specific Problem

Today's only option for Caruca-style specification mining is v1's hand-written, hard-to-extend pipeline (6,520 LOC published — of which roughly 3,130 is LLM-generated specification data, so the hand-written pipeline is about **3,456 LOC**; measured 2026-08-29, see `memory/caruca_v1_loc_baseline.md`). Any claim about replacing hand-written logic must use 3,456, not 6,520; the LLM is used only for syntax-spec inference, and even that step is currently broken against modern DSPy). This leads to slow iteration and no existing measurement of what an LLM-first approach would cost or how well it would work.

The problem is sharper than just v1's maintainability: v1's pipeline is fundamentally scoped to commands with well-documented man pages and a testable, non-destructive execution model. It doesn't generalize to arbitrary executables: binaries with sparse or no documentation, tools whose behavior depends on binary format inspection rather than text parsing, or commands (like `rm`) that are destructive enough to require a secure sandbox before any invocation permutation can even be run. Before an LLM-based approach can be

judged fairly against v1, the evaluation itself needs infrastructure v1 never had to build: a safe execution environment for destructive commands, and a plan for what auxiliary tooling an LLM would need to handle commands outside the man-page-documented, non-destructive sweet spot v1 was designed for.

There's a second, more basic asymmetry worth naming: v1's behavior is fixed by hand-written code, so it has no configuration surface at all outside its one LLM call. Once v2 replaces hardcoded logic with an NLP-based component, model choice, decoding parameters (e.g. temperature), and output-format constraints (e.g. structured/JSON output) stop being incidental setup and become part of the method itself, directly shaping accuracy, consistency, and cost. Selecting and documenting them deliberately is core scope for v2, not optional tuning bolted on afterward, and characterizing how much they matter is itself new evaluative territory nobody has covered for this task.

Stakeholders & Consumers

Research Stakeholders

- Tiran Dagan — implementer
- Prof. Michael Greenberg (Stevens) — PI/advisor for Caruca; primary reviewer of Caruca-specific results
- Prof. William Eiers (Stevens) — PhD advisor; methodology sounding board, not Caruca-specific

White Paper Co-Authors

Prospective co-authors on the paper resubmission, and evaluators of the approach in the meantime:

- Evangelos Lamprou (Brown)
- Seong-Heon Jung (NYU)
- Mayank Keoliya (UPenn)
- Lukas Lazarek (Brown)
- Konstantinos Kallas (UCLA)
- Nikos Vasilakis (Brown)

Downstream Spec Consumers

Systems the output must satisfy, same as v1's, per the paper's §7:

- PaSh, POSH, ShellCheck, Shseer — each expects its own annotation shape

Operators

- Whoever runs the tooling day to day — currently just Tiran via CLI; revisit if a web GUI enters scope

Evaluation Criteria & Success Definition

1. **Correctness/fidelity** — how closely produced specifications/annotations match v1 and ground truth (reusing v1's §7.2 `cmp_specs.py` methodology)
2. **Cost/performance** — wall-clock, tokens, \$ per command, captured for both v1's LLM step and the new system, so we can say whether v2 matches, exceeds, or trails v1 on cost/speed
3. **Coverage** — command behaviors, flags, and configurations handled, including how far v2 extends beyond v1's man-page-documented, non-destructive sweet spot
4. **Consistency/reproducibility** — variance across repeated runs (LLM nondeterminism is itself a reportable result)
5. **Percent-change roll-up** — an explicit v1-vs-v2 delta computed per dimension above (e.g. "+/-N% on Q2 exact-match," "Nx cost"), giving the comparison one clear headline number per dimension
6. **Reduction in hand-encoded logic** — how much procedural code is replaced for equal-or-better output on dimensions 1-4

No concrete target numbers yet; this is deliberately qualitative at this stage.

MVP Components (Flexible, provisional)

- **Baseline instrumentation** (extends v1)
 - Add token/cost/wall-clock/model-ID telemetry to v1's existing LLM step, so it's a fair comparison point
- **Naive-LLM baseline** (new, per Eiers' guidance)
 - CLI tool: command binary + docs in, specification in a downstream-consumable format out
 - Few-shot primed with concrete worked examples
- **Secure sandbox environment** (new capability v2 adds beyond v1's model)
 - Isolated execution for arbitrary command permutations, including destructive ones (e.g. `rm`)
 - Must preserve strace/ptrace-level syscall visibility (hard constraint from the tracer)
 - Candidate approaches (research done, not yet decided): extend v1's Docker/overlayfs backends for bulk CLI sweeps; Firecracker microVMs on the CloudLab host for stronger isolation on destructive commands (real guest kernel keeps `strace` working natively); gVisor as a possible middle ground pending a ptrace-compatibility spike
 - One sandbox service, two front doors: the CLI calls it directly, the future web harness calls the same service over an API
- **LLM tool-augmentation layer** (new capability v2 adds beyond v1's model)
 - Starter set: binary/format inspection (e.g. binhex-style dumps), static/code analysis, help/usage-string extraction, for executables outside the man-page/text-documentation sweet spot
 - Likely to grow (not yet committed): web search for documentation, GitHub connectivity for source/issue lookup, as specific target commands demand them
- **Evaluation harness**
 - Reuses v1's `eval/cmp_specs.py` methodology and `benchmarks/annotations/` ground truth
 - Produces the three-way comparison (v1 vs. naive-LLM vs. ground truth) across the 6 evaluation dimensions
 - Computes the percent-change roll-up (dimension 5) as an explicit output, not just raw numbers
 - Output should be usable directly in a paper (tables/figures/methodology description), given the resubmission objective
- **Web GUI** (*optional, not yet decided*)

- For running/visualizing comparisons interactively instead of via CLI/scripts
- Stretch idea: in-browser CPU emulation (v86-style) as a fully client-side “try it live” demo mode

[Add or adjust components as system design firms up]

Key Usage Scenarios

Researcher (Tiran)

1. **Run a baseline comparison** *As the researcher, I want to run the naive-LLM baseline against a command and get back a spec plus token/cost/latency numbers, so that I can add a row to the comparison table without reconstructing measurements after the fact.*
2. **Diff against ground truth** *As the researcher, I want to diff the naive-LLM’s spec against v1’s ground truth using the existing `cmp_specs.py` methodology, so that correctness numbers are directly comparable to the paper’s Q2 results.*
3. **Safely sweep destructive commands** *As the researcher, I want to run invocation permutations of destructive commands (e.g. `rm`) inside the secure sandbox, so that coverage testing isn’t limited to the non-destructive commands v1’s model could safely handle.*
4. **Track v1-to-v2 deltas as they happen** *As the researcher, I want every methodological deviation between v1 and v2 logged as it’s made, so that the paper resubmission’s “what changed and why” section can be written from records, not reconstructed from memory.*

Advisor/Reviewer (Greenberg, Eiers, co-authors)

1. **See the three-way comparison** *As a reviewer, I want the v1 / naive-LLM / ground-truth comparison summarized per dimension with the percent-change roll-up, so that I can judge how far the LLM-based approach extends what v1 could do, and whether it’s strong enough to support a resubmission.*

Author Team (Tiran + co-authors, for resubmission)

1. **Paper-ready output** *As the author team, we want the evaluation harness’s results in a form usable directly in a paper (tables/figures/methodology description), so that drafting the resubmission doesn’t require re-deriving results from scratch.*

System/Background

1. **Automatic telemetry capture** — When a baseline or naive-LLM run completes, then its telemetry (tokens, cost, wall-clock, model ID, seed) is recorded alongside the output, so no measurement has to be reconstructed later.
2. **Sandbox teardown** — When a sandboxed execution finishes (or crashes), then the isolated environment is torn down automatically, so no destructive side effect can leak to the host regardless of outcome.

Future / Stretch Directions (*not committed*)

- **Agentic full-pipeline rebuild** — replacing configuration generation, tracing orchestration, and specification derivation with Claude Code/Agent-SDK reasoning (Tiran's original proposal; deferred per Eiers until the naive baseline + evaluation harness exist)
- **Web GUI/dashboard** — for running and visualizing comparisons interactively, if the CLI-first approach turns out to be limiting
- **In-browser CPU emulation demo mode** — a fully client-side, zero-server-trust way to try a single command live in the web GUI (v86/TinyEMU-style), once the web GUI itself is in scope
- **Expanded tool-augmentation** — web search for documentation, GitHub connectivity for source/issue lookup, added as specific target commands demand them
- **Expanded command coverage** — beyond whatever initial command set the naive baseline targets first